



Astronomical Image Convolution

Christian Harris

Convolution



- Intentional blurring of an image to highlight specific features



Large Scale Features

Convolution

- Intentional blurring of an image to highlight specific features




Small Scale Features

Convolution

Combining neighboring pixels to produce a single value representative of that region

- *Kernel* determines how neighboring values are related



Kernel Types


Boxcar

1	1	1
1	1	1
1	1	1



Laplacian

-1	-1	-1
-1	8	-1
-1	-1	-1



Gaussian

0.21	0.29	0.21
0.29	0.38	0.29
0.21	0.29	0.21

Computer Specs

- Ubuntu 20.04
- Intel(R) Core(TM) i7-7700HQ CPU @ 2.80GHz
- 4 physical cores
- 8 threads

OpenCilk

Compiling and running Cilk

Compiling a Cilk program is similar to compiling an ordinary C or C++ program. To compile a Cilk program using Tapir/LLVM, add the `-fcilkplus` flag to the `clang` or `clang++`

Porting Cilk Plus code to OpenCilk


To port a Cilk Plus program to OpenCilk, once all uses of unsupported features have been updated, make the following changes to your build process:


- When compiling the program, replace any uses of `-fcilkplus` with `-fopencilk`.
- When linking the program, replace any uses of `-lcilkrt` with `-fopencilk`.

- Installed from source
 - ~ 45 minutes
- First attempt crashed at 12%
 - Documented mismatch in gcc and g++ compiler version
- Compiler issues
 - Regular clang++ didn't contain OpenCilk functionality
 - Had to call clang++ from the OpenCilk build directory
- Documentation mismatch
 - `-fopencilk`

Overview

- Scanning over an image
- At each pixel, apply the kernel
 - Multiply each kernel element with the associated image pixel
 - Sum
 - Assign value to new image pixel

$$C[x, y] = \sum_{i=1}^{(2V+1)} \sum_{j=1}^{(2W+1)} K[i, j] A[(x - V - 1 + i), (y - W - 1 + j)]$$




Single Pixel Convolution

```
int convolvePixel(int x,int y, Matrix& image, Matrix& kernel) {
    int kernellength = kernel.getRow();
    int V = (kernellength - 1)/2;

    int counter = 0;
    int convolutionVal = 0;
    // Center the Kernel above the Image pixel location matching the NewImage pixel location
    for (int i = 0; i < kernellength; i++) {

        for (int j = 0; j < kernellength; j++) {
            // Current Image pixel being applied to ConvolutionVal
            int imagX = x-V+i;
            int imagY = y-V+j;

            // Exclude indices outside Image (edge detection)
            if ( ((imagX >= 0) && (imagX < image.getRow())) &&
                ((imagY >= 0) && (imagY < image.getCol())) )
            {
                counter += kernel.get(i,j); // Track Kernel values used (flux conservation)
                convolutionVal += kernel.get(i,j) * image.get(imagX, imagY);
            }
        }
    }
    return convolutionVal / counter;
}
```

Sequential Convolution

```
void sequentialConvolve(int startX, int endX, int startY, int endY, Matrix& convolvedImage, Matrix& originalImage, Matrix& kernel) {  
    for (int i = startX; i < endX; i++) {  
        for (int j = startY; j < endY; j++) {  
            convolvedImage.set(i,j, convolvePixel(i,j,originalImage,kernel));  
        }  
    }  
}
```

- Takes a region to convolve as input
- Calls convolvePixel() for each pixel in the region

Parallel Convolution

```
void parallelConvolve(int startX, int startY, int xSize, int ySize, Matrix& convolvedImage, Matrix& originalImage, Matrix& kernel, int stopX, int stopY) {  
    if ((xSize <= stopX) && (ySize <= stopY)) {  
        sequentialConvolve(startX, startX+xSize, startY, startY+ySize, convolvedImage, originalImage, kernel);  
    }  
    else {  
        cilk_spawn  
            parallelConvolve(startX, startY, xSize/2, ySize/2, convolvedImage, originalImage, kernel, stopX, stopY);  
        cilk_spawn  
            parallelConvolve(startX, startY+ySize/2, xSize/2, ySize/2, convolvedImage, originalImage, kernel, stopX, stopY);  
        cilk_spawn  
            parallelConvolve(startX+xSize/2, startY, xSize/2, ySize/2, convolvedImage, originalImage, kernel, stopX, stopY);  
        cilk_spawn  
            parallelConvolve(startX+xSize/2, startY+ySize/2, xSize/2, ySize/2, convolvedImage, originalImage, kernel, stopX, stopY);  
        cilk_sync;  
    }  
}
```

- Divide and Conquer subproblems of size $n/2$
- Assumes images are not square
- Stopping size is a parameter

Testing Objectives

- 3x3 Boxcar kernel (all ones)
- Typical detector sizes
 - 1024x1024
 - 2048x2048
 - 4096x4096
 - 8196x8196
- Typical pixel values
 - 16-bit unsigned int -> [0, 65535]
 - Measured in analog digital units

- Varying stopping sizes
 - 1
 - 5% original size
 - 10% original size
 - 20% original size

Testing

```
int n = 8192;
dataSource.open("Image_8192.txt");
Matrix image = Matrix(dataSource, n);
dataSource.close();


Matrix newPimage = Matrix(n);
int stopSize = 1;

auto start = std::chrono::high_resolution_clock::now();
parallelConvolve(0, 0, newPimage.getRow(), newPimage.getCol(), newPimage, image, kernel, stopSize, stopSize);
auto stop = std::chrono::high_resolution_clock::now();


auto duration = std::chrono::duration_cast<std::chrono::microseconds>(stop - start);
std::cout << "Parallel:   " << duration.count() << " microseconds\n";
```

- Run multiple times changing CILK_NWORKERS


1024



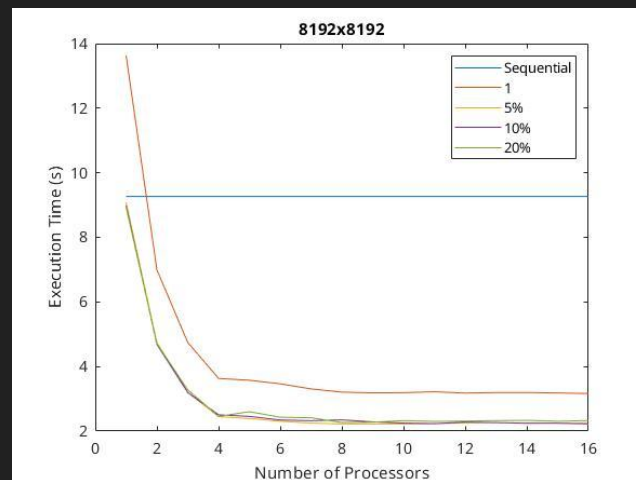
2048



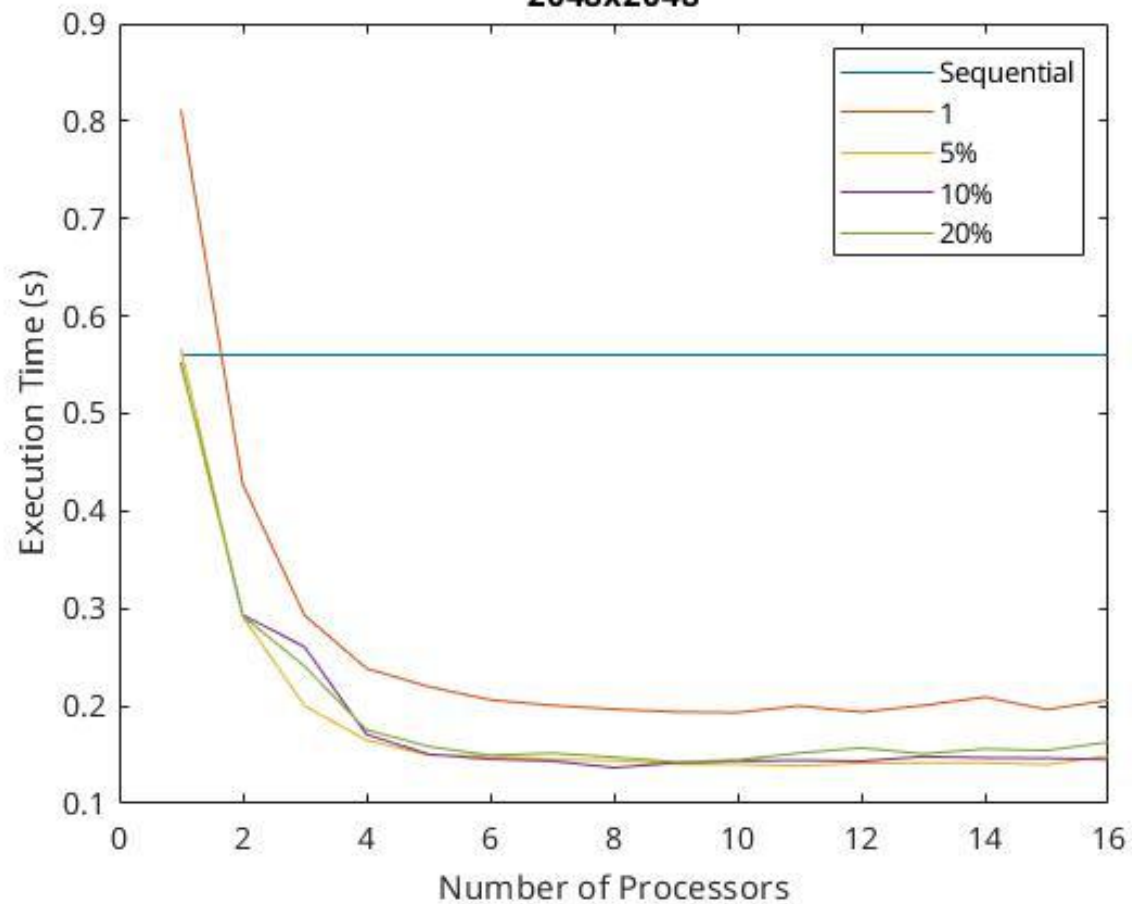
4096



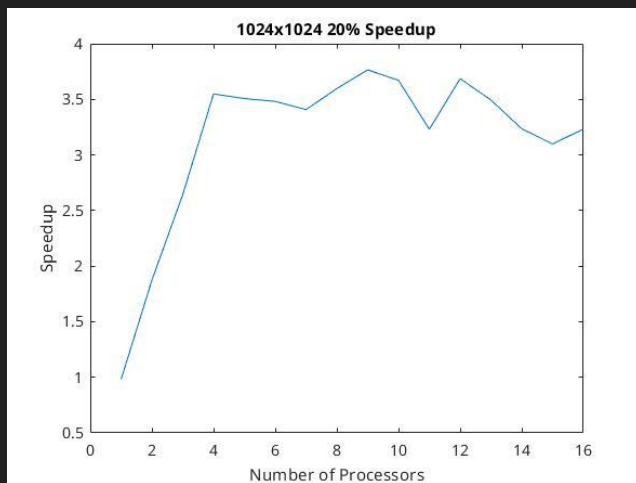
8192



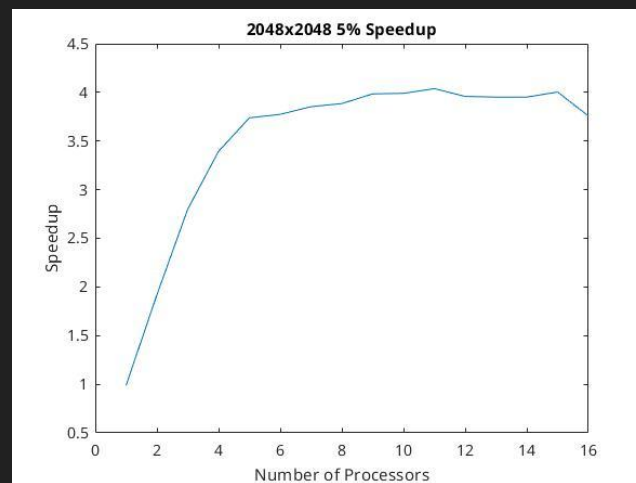
2048x2048




1024



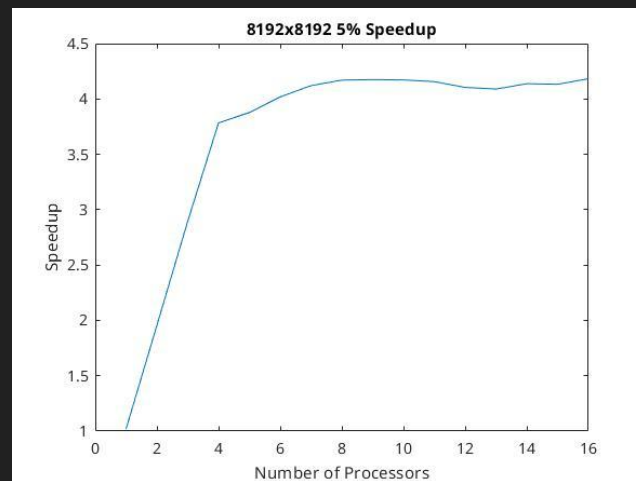
2048

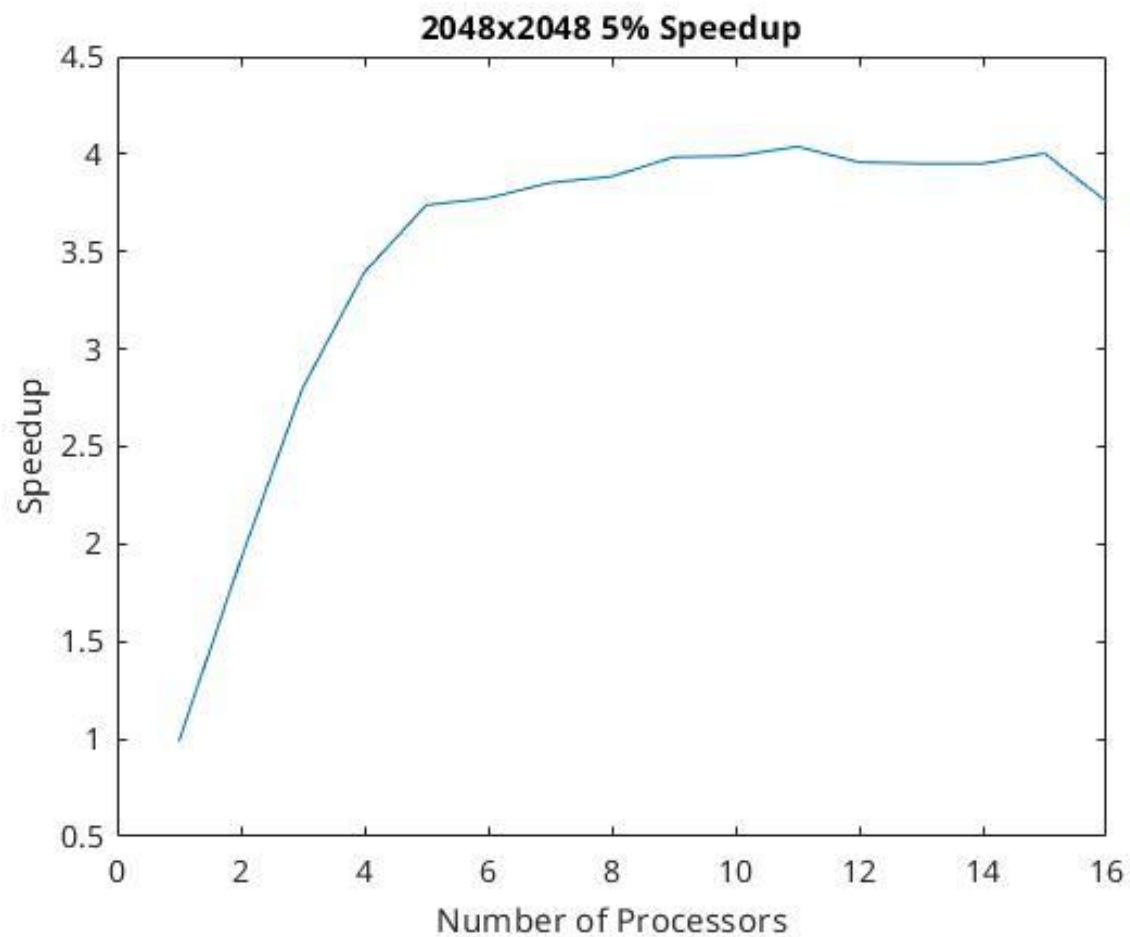


4096



8192





Analysis

- Significant decrease in execution time by not subdividing problem to size 1
- BUT, why no variation as stopping size increases?
 - Larger subproblems done sequentially, yet same execution time
 - Bug in code?
- Plateau at 4 processors
 - Gains essentially stop after physical cores exhausted

Thank You

Questions?